cloudera®

A [Incomplete] Data Tools Landscape [for Hackers] in 2015

Wes McKinney @wesmckinn

Data^3 Meeting — Minneapolis, MN

This talk

- A partial look at different languages and tools
- Limiting scope to either:
  - Permissively licensed open source software, e.g. Apache-licensed (OSS)
  - Non-dual-licensed copyleft OSS (e.g. GPL)
  - i.e. “do you [the community] have any incentive to create patches?”
- Some trends (that I see, anyway)
- Challenges and opportunities

Who am I?

- Python data firestarter
- Financial analytics in R / Python starting 2007
- pandas project born of frustration in 2008
- 2010-2012
  - Hiatus from gainful employment
  - Make pandas ready for primetime
  - Write "Python for Data Analysis"

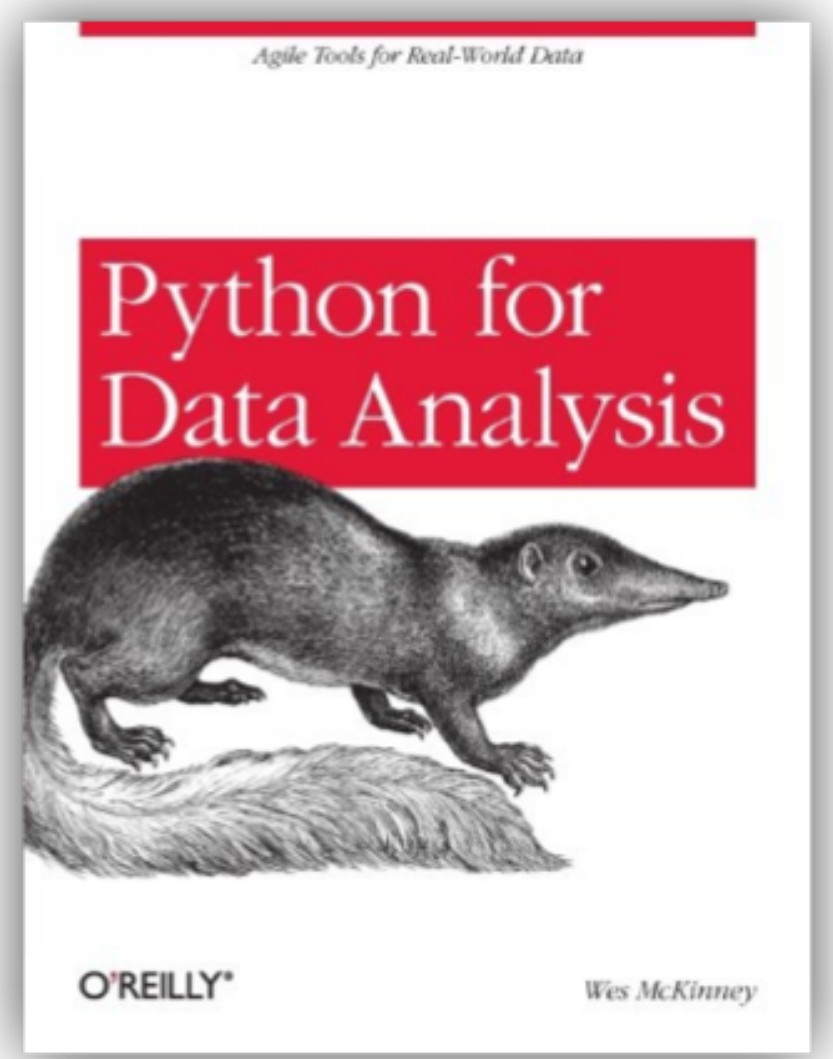

cloudera

Who am I? (cont'd)

- 2013-2014: Co-founder/CEO of DataPad (analytics startup, with early pandas collaborator Chang She)
- Late 2014: DataPad team joins Cloudera
- Now: backend systems and all-things-Python @ Cloudera

SQL: Still a lingua franca

- "SQL: the Fortran of Analytics"
- Often a concise, declarative way to express data transforms, analytics, etc.
- Relatively easy to parse, analyze
- SQL recently has seen resurgence with focus on interactive-speed SQL engines, especially on top of HDFS/Hadoop
- Relevant and impactful features (e.g. JSON support) still arriving in established RDBMS like PostgreSQL

Historical Python Context

- Scientific / HPC computing focus in 1990s, 2000s
  - Python web community developed in parallel, matured faster!
- NumPy became community standard in 2005, born from Numeric + Numarray
- Pyrex, later Cython, easier C / C++ wrapping
- f2py: easy Fortran wrapping
- Anaconda distribution
  - Finally solving Python deployment for all

Essential Python stack

- NumPy: low-level array processing
- SciPy: essential computational algos
- pandas: data wrangling
- scikit-learn: machine learning
- matplotlib (+ add-ons, like seaborn): visualization
- numba: numeric hotspot LLVM compiler
- Domain-specific toolkits: nltk, scikit-image, statsmodels, Theano, PyCUDA/ PyOpenCL and many others

pandas

- A Pythonic take on the classic R “data frame” data structure
- Critical piece to make the Python stack useful in everyday work
- Added axis metadata / labeling for representing multidimensional data
- Focus on easy data wrangling, IO, plotting, and basic analytics

cloudera

Jeff Reback's "pandas as PyData middleware" diagram

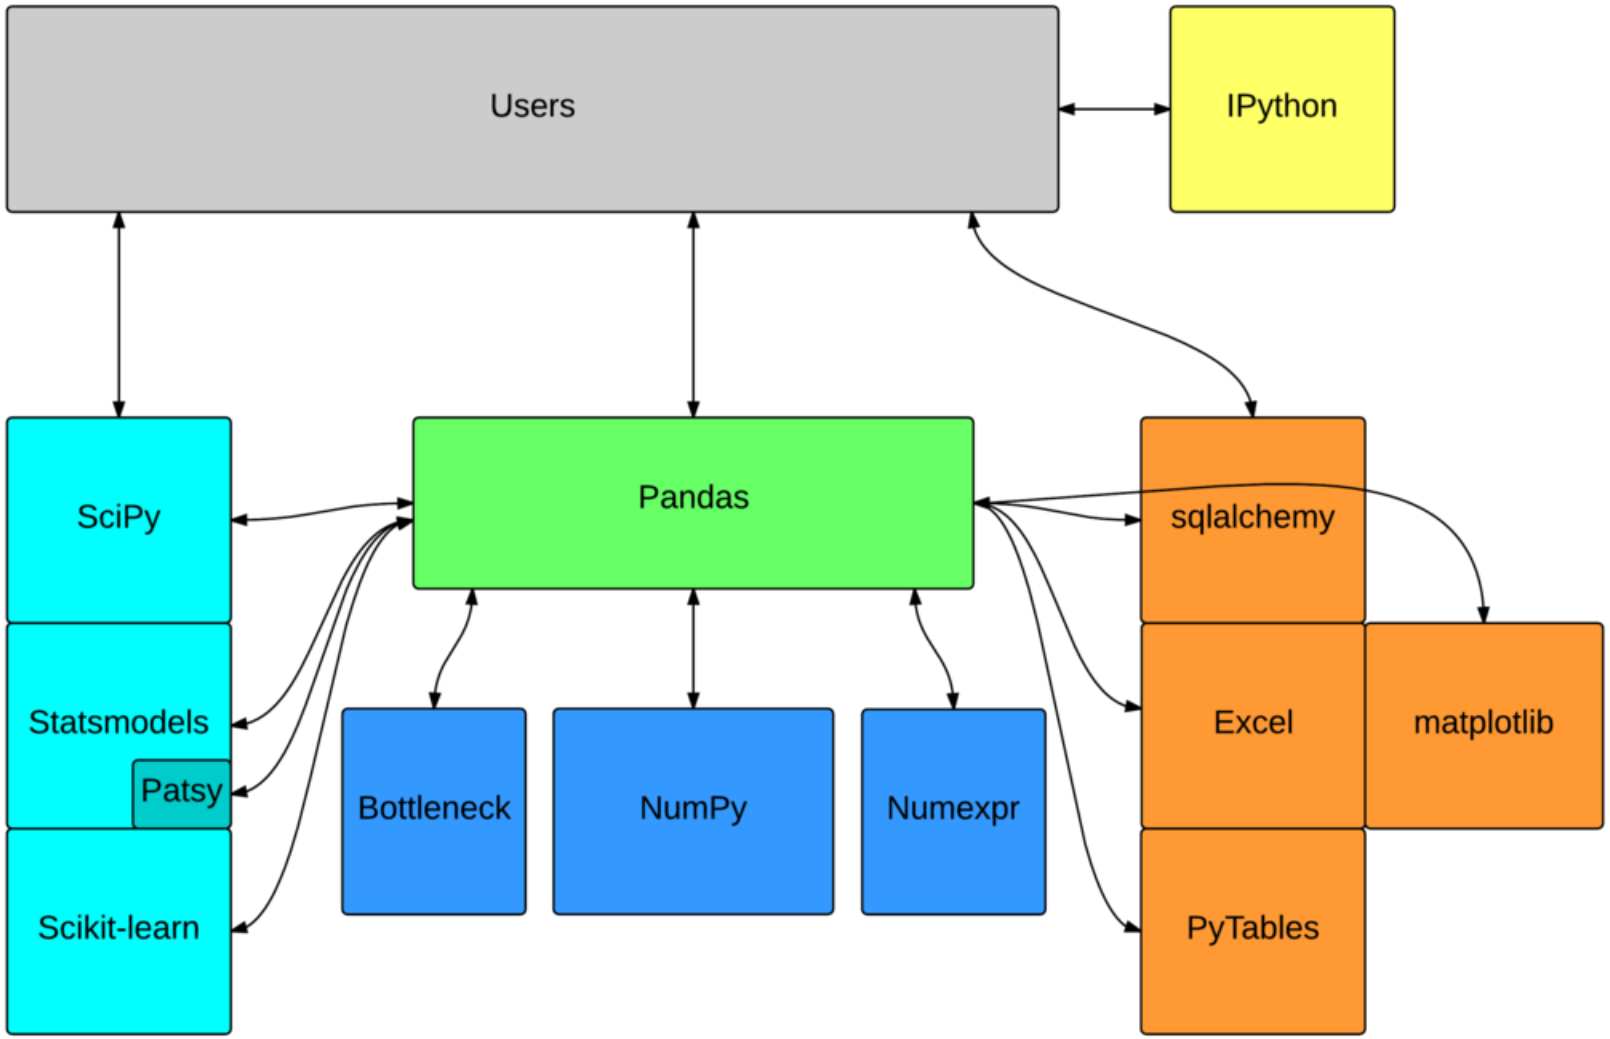

cloudera

Newer / Up-and-coming Python projects

- Bokeh: interactive / reactive visualization for the web
- Blaze: uniform data expression API
- Odo: easy data migration

R Project

- Trusted base of statistics libraries
  - Latest and greatest stats research often hits R first
- RStudio
- The "Hadley stack”
  - Visualization: ggplot2 (static) and ggvis (interactive)
  - Data Wrangling: dplyr
  - legacy: plyr / reshape2

cloudera

dplyr

- Started late 2012 by Hadley Wickham, supported by RStudio
- Composable / chainable analytics and data wrangling expressions
- In-memory and SQL backends
- Has attracted folks back to R from Python in a lot of cases

Some other great R stuff

- shiny: interactive web apps in R
- Rcpp
- data.table
- xts

IPython

- IPython started out as a better interactive Python
- Grew to include web-based computational notebook, GUI console, and other components
  - (Google even integrated into Google Drive!)
- IPython Notebook architecture enabled “kernel” processes to be written in nearly any language (even bash!)
- How to build community beyond Python?

Enter Jupyter

- http://jupyter.org
- Breaking out notebook machinery into a standalone non-Python-specific project
- Enable project components to evolve at own pace, without large monolithic releases
- JupyterHub: upcoming multi-user notebook server

cloudera

A few words about Hadoop + Big Data

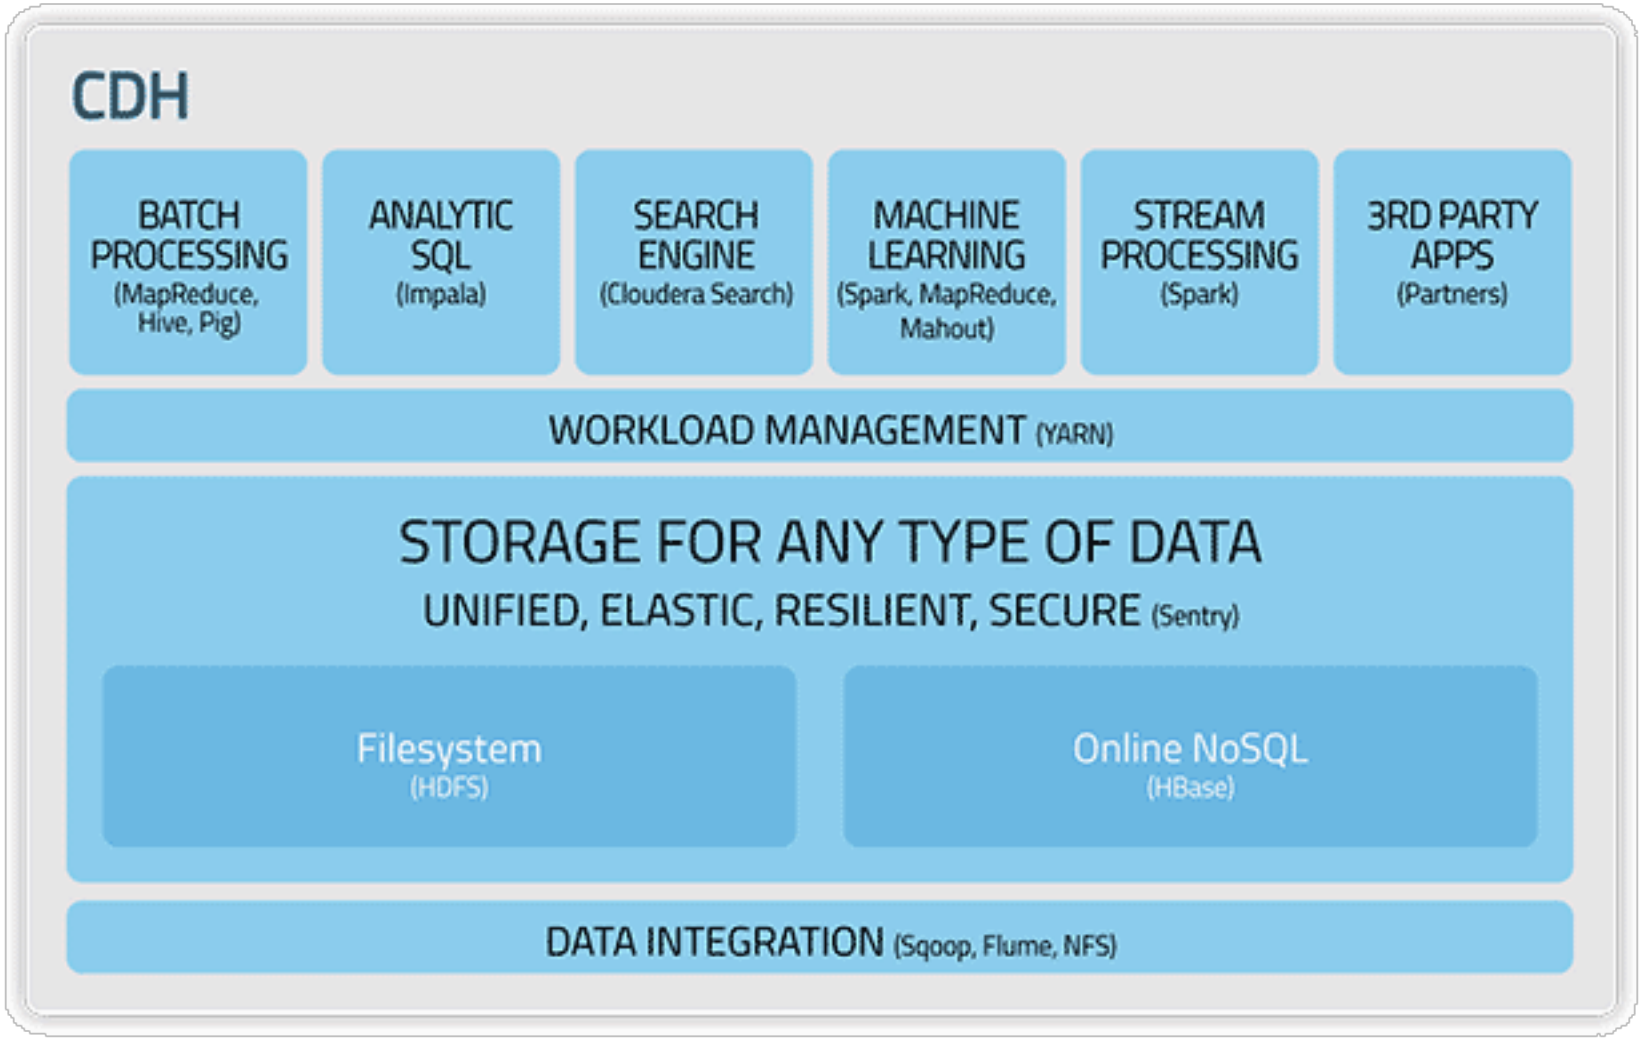

cloudera

Apache Spark

- Originated from Berkeley AMPLab
- General purpose distributed memory-centric data processing framework
- Official APIs: Scala, Java, Python

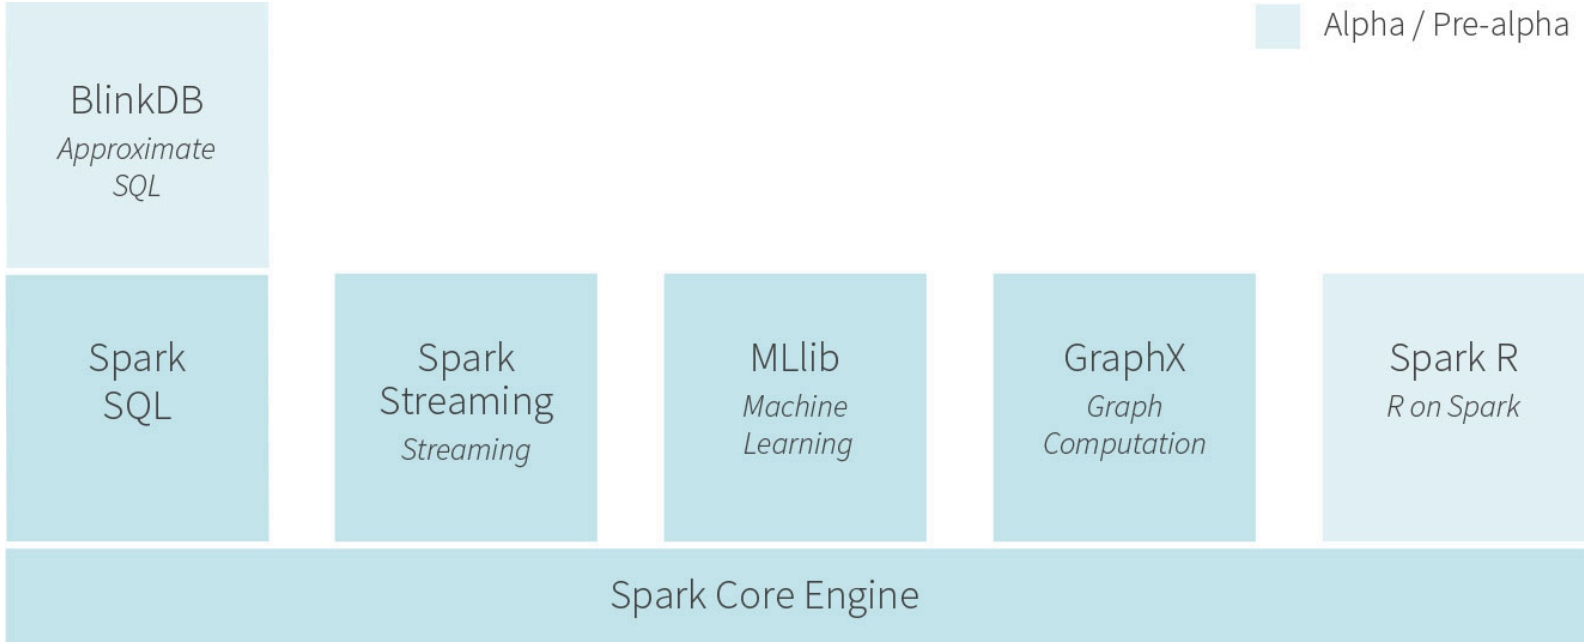

Source: databricks.com

cloudera

Spark 1.3: DataFrames!

- R/pandas-inspired API for tabular data manipulation in Scala, Python, etc.
- Logical operation graphs rewritten internally in more efficient form
- Good interop with Spark SQL
- Some interoperability with pandas
- Will help close the semantic gap between Spark and R/Python

Some problems in need of solving

- A Shiny-like quick-and-dirty data app development framework for Python
- IPython/Jupyter notebook collaboration
- A community-standard, Apache-licensed C/C++ data frame library with best-in-class performance
- Ubiquitous support for emerging analytical on-disk storage standards like Parquet

Other interesting stuff to look at

- Torch7 / LuaJIT: high performance ML / deep learning on GPUs
  - Facebook AI group open sourced several ML modules
- Apache Flink
  - Up-and-coming Scala-based data processing framework
  - Some overlap with Spark use cases

Some other interesting industry trends

- Microsoft
  - Acquired Revolution Analytics, leading commercial R vendor
  - Launched Azure ML: R, Python, and more on Azure cloud
- Dato (fka GraphLab)
  - faster, more scalable machine learning, with Python interface (Paid commercial product, free for non-commercial/academic use)
  - Largest-ever VC investment in a data tools company betting big on Python
- Databricks
  - Offering cloud Spark-notebook-as-a-service

cloudera

cloudera

Thank you

@wesmckinn

